



Shared-memory Programming with OpenMP

Advanced Computing Laboratory
Costa Rica National High Technology Center

Diego Jiménez Vargas
djimenez@cenat.ac.cr

2018

Contents

1 Introduction

- Parallel Architectures
- Parallel Computing
- Directive-based Programming
- Warm-up Exercise

2 OpenMP

- Standard
- Parallel Directive and Variables
- Loop Directives
- Additional Directives
- Shared-memory Systems

3 Exercises

Introduction

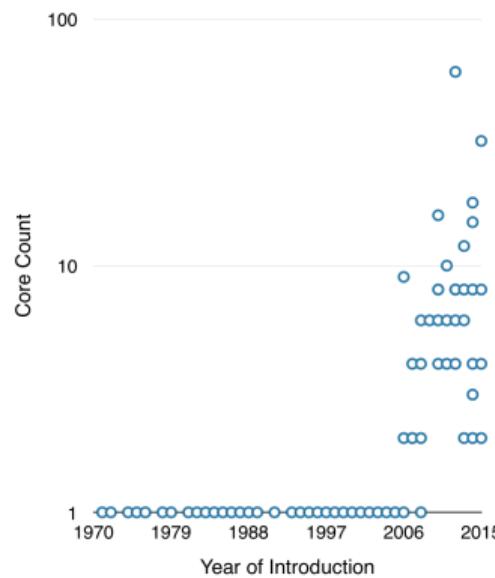
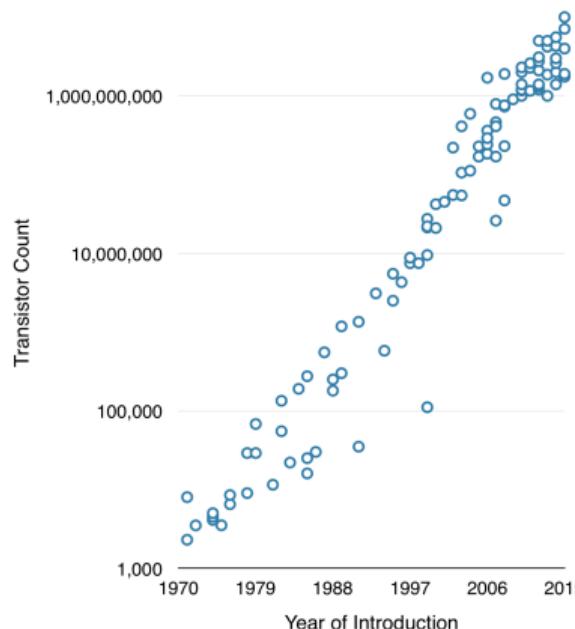
Parallel Architecture

It's everywhere



Moore's Law

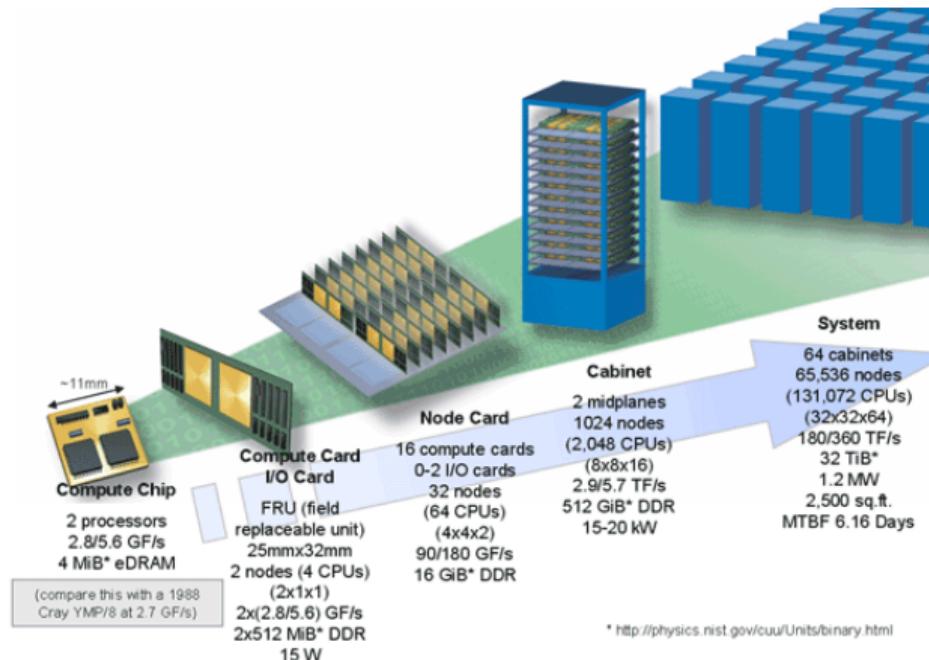
Transistor density increases exponentially



Source: https://en.wikipedia.org/wiki/Transistor_count

Supercomputer

A hierarchical ensemble of many parts



Top 500

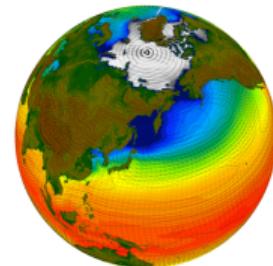
Fastest supercomputers on Earth

Rank	System	Cores	Rmax (TFlop/s)	Rpeak (TFlop/s)	Power (kW)
1	Sunway TaihuLight - Sunway MPP, Sunway SW26010 260C 1.45GHz, Sunway , NRCPC National Supercomputing Center in Wuxi China	10,649,600	93,014.6	125,435.9	15,371
2	Tianhe-2 [MilkyWay-2] - TH-IVB-FEP Cluster, Intel Xeon E5-2692 12C 2.200GHz, TH Express-2, Intel Xeon Phi 31S1P , NUDT National Super Computer Center in Guangzhou China	3,120,000	33,862.7	54,902.4	17,808
3	Piz Daint - Cray XC50, Xeon E5-2690v3 12C 2.6GHz, Aries interconnect , NVIDIA Tesla P100 , Cray Inc. Swiss National Supercomputing Centre (CSCS) Switzerland	361,760	19,590.0	25,326.3	2,272
4	Gyoukou - ZettaScaler-2.2 HPC system, Xeon D-1571 16C 1.3GHz, Infiniband EDR, PEZY-SC2 700Mhz , ExaScaler Japan Agency for Marine-Earth Science and Technology Japan	19,860,000	19,135.8	28,192.0	1,350
5	Titan - Cray XK7, Opteron 6274 16C 2.200GHz, Cray Gemini interconnect, NVIDIA K20x , Cray Inc. DOE/SC/Oak Ridge National Laboratory United States	560,640	17,590.0	27,112.5	8,209

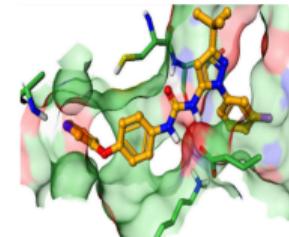
Source: <https://www.top500.org/> (November 2017)

Advanced Computing

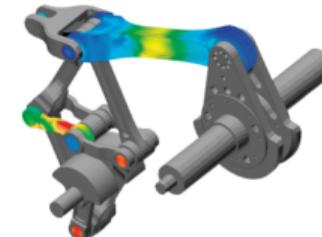
For everyone



Climatology



Medicine



Materials



Economics



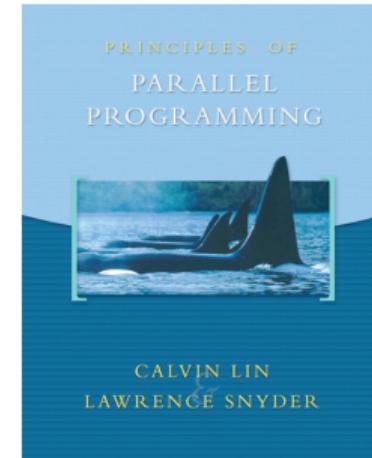
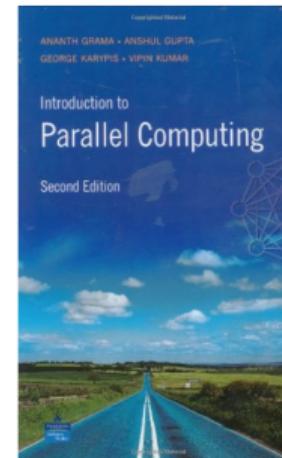
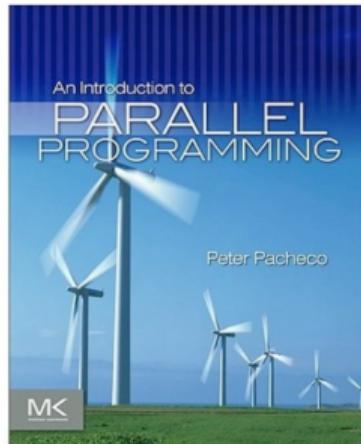
Cosmology



Entertainment

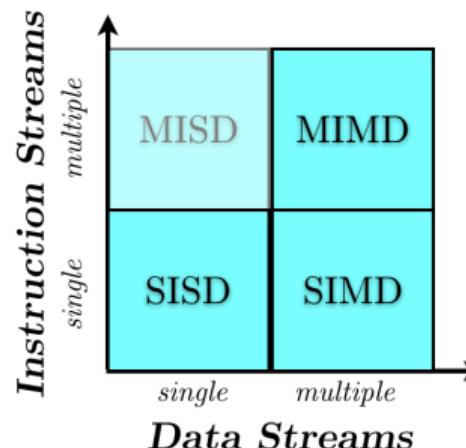
Parallel Computing Books

Do they look similar?



Flynn's Taxonomy

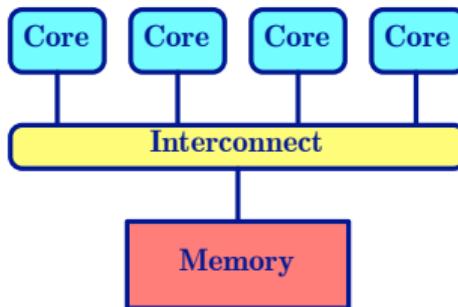
A classification of computer systems



- SISD (Single Instruction, Single Data): traditional von-Neumann model of single-processor computers
- SIMD (Single Instruction, Multiple Data): Vector processors and accelerators
- MISD (Multiple Instruction, Single Data): no system fits this description
- MIMD (Multiple Instruction, Multiple Data): standard computer clusters:
 - **Shared-memory systems**
 - Distributed-memory systems

Shared-memory Systems

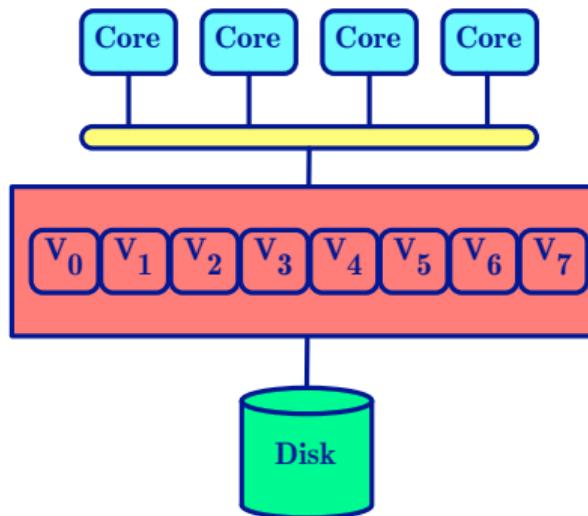
Single address space



- Communication is implicit: write and read operations on shared variables
- Simple programming model: no data distribution among processors
- Limited scalability (memory contention)
- Multicore processors fit this description
- Symmetric Multiprocessor (SMP): memory access latency is the same for all processors, also called Uniform Memory Access (UMA)
- Non-Uniform Memory Access (NUMA) systems are more scalable

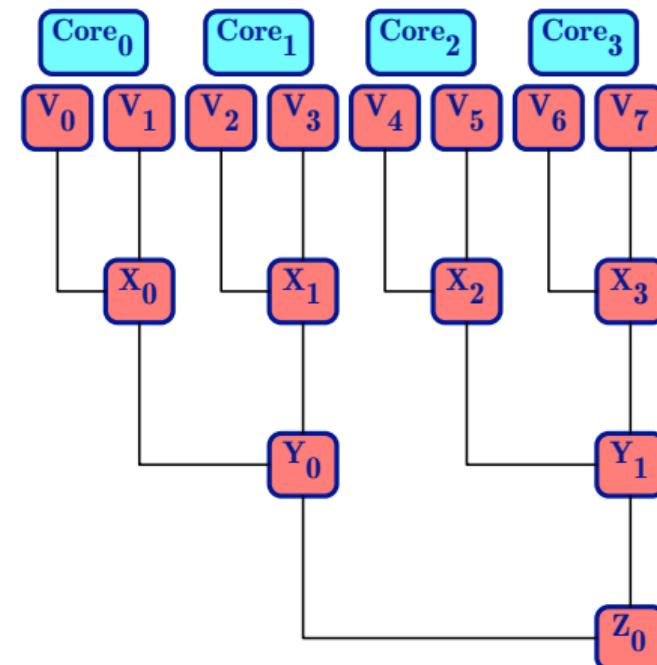
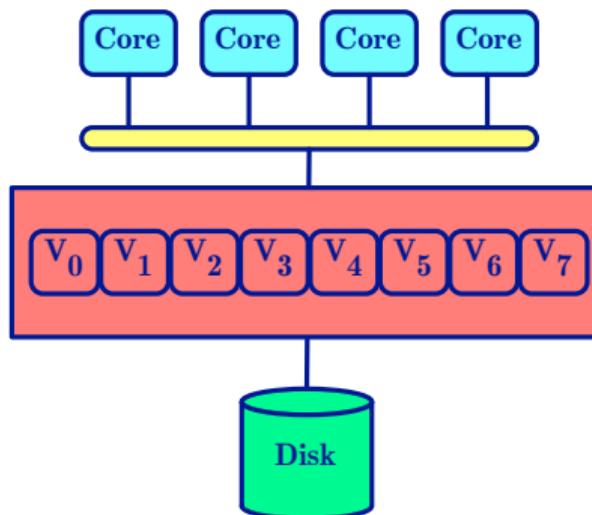
Parallel Computing

An example: parallel sum



Parallel Computing

An example: parallel sum



Speedup and Efficiency

Performance metrics

Definition (Speedup)

The ratio of the best sequential time over the parallel time with p cores:

$$S_p = \frac{T_1}{T_p}$$

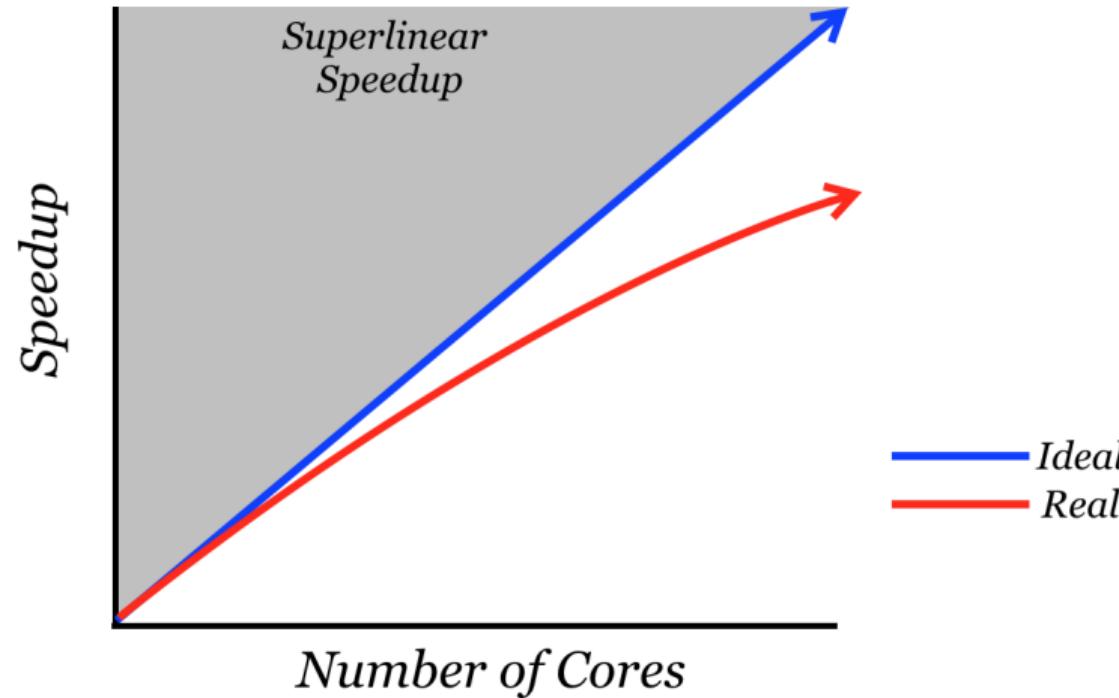
Definition (Efficiency)

The average utilization of a processor in the parallel system:

$$E_p = \frac{S_p}{p}$$

Speedup Curve

Increases with more cores until it saturates



Amdahl's Law

Bound to speedup

Theorem (Amdahl's Law)

Assume a fraction s of the work for a given problem is serial, with $0 \leq s \leq 1$, while the remaining portion $1 - s$ is p -fold parallel. Therefore,

$$T_p = sT_1 + (1 - s)\frac{T_1}{p}$$

$$S_p = \frac{T_1}{T_p} = \frac{p}{sp + (1 - s)} \quad E_p = \frac{T_1}{pT_p} = \frac{1}{sp + (1 - s)}$$

Hence, $S_p \rightarrow \frac{1}{s}$ and $E_p \rightarrow 0$ as $p \rightarrow \infty$.

For instance, if the serial fraction exceeds 1 percent, then speedup can never exceed 100 for any p .

Scalability

Problem size importance

Definition (Scalable)

A technology is scalable if it can handle ever-increasing problem sizes.

Definition (Strong scalability)

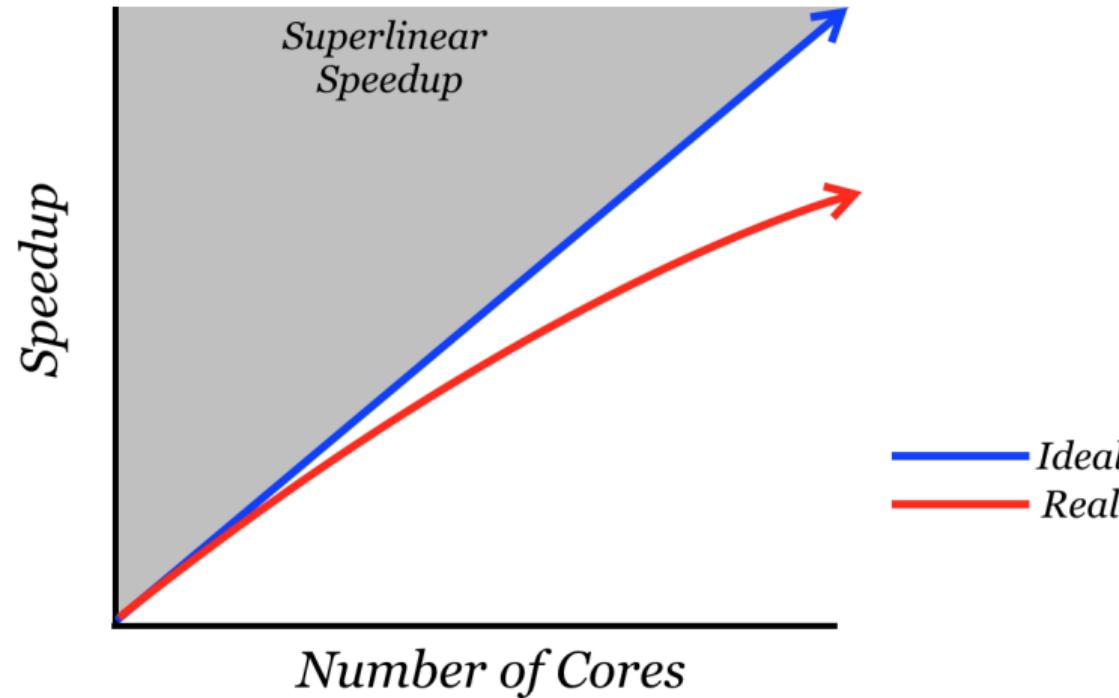
If when we increase the number of processes/threads, we can keep the efficiency fixed *without* increasing the problem size, we say a program is *strongly scalable* (Amdahl's law).

Definition (Weak scalability)

If we can keep the efficiency fixed by increasing the problem size at the same rate as we increase the number of processes/threads, we say a program is *weakly scalable* (Gustafson's law).

Speedup Curve

Increases with more cores until it saturates



Exercises

- ① What is the maximum expected speedup of a system with 1024 processors?
- ② What is the maximum expected efficiency of a system?
- ③ What is the efficiency of a 256-processor system if speedup is 64?
- ④ What is the speedup of a system with efficiency 0.5 and 32 processors?
- ⑤ What is the maximum speedup achievable on a parallel system if the serial portion of the program represents 20% of sequential execution time?

Compilers

Translate high-level code into low-level representation

Fortran

```
1 do i=1,n  
2 C(i) = A(i) + B(i)  
3 enddo
```

Assembly

```
1 Label1:  
2     mov al,[si]  
3     add al,[di]  
4     mov [bx],al  
5     inc si  
6     inc di  
7     inc bx  
8     dec cl  
9     jnz Label1
```

Compilers

Translate high-level code into low-level representation

Fortran

```

1 do i=1,n
2 C(i) = A(i) + B(i)
3 enddo

```

Assembly

```

1 Label1:
2     mov al,[si]
3     add al,[di]
4     mov [bx],al
5     inc si
6     inc di
7     inc bx
8     dec cl
9     jnz Label1

```

Compiler optimizations (gcc -O3)

```

-fauto-inc-dec
-fbranch-count-rcs
-fcombine-stack-adjustments
-fcompare-elim
-fcompare-registers
-fdce
-fdefer-pop
-fdelayed-branch
-fdse
-fforward-propagate
-fguess-branch-probability
-fif-conversion
-fif-conversion
-finline-functions-called-once
-fipa-pure-const
-fipa-profile
-fipa-reference
-fmerge-constants
-fmove-loop-invariants
-fshrink-wrap
-fsplit-wide-types
-ftree-bit-csp
-ftree-cpp
-ftree-phipt
-ftree-ch
-ftree-copy-prop
-ftree-copyrename
-ftree-dce
-ftree-dominator-opts
-ftree-dse
-ftree-formprop
-ftree-fre
-ftree-phirop
-ftree-sink
-ftree-slaer
-ftree-sra
-ftree-ptx
-ftree-ter
-funit-at-a-time
-ftree-built-in-call-dce
-ftree-switch-conversion
-ftree-tail-merge
-ftree-pre
-ftree-vrp
-fica-ra
-fthread-jumps
-falign-functions
-falign-jumps
-falign-loops
-falign-labels
-fcaller-saves
-fcrossjumping
-fcase-follow-jumps
-fcase-skip-blocks
-fdelete-null-pointer-checks
-fdevirtualize
-fdevirtualize-speculatively
-fexpensive-optimizations
-fcase
-fcase-lm
-fhoist-adjacent-loads
-finline-small-functions
-findirect-inlining
-fipa-cp
-fipa-cp-alignment
-fipa-sra
-fipa-icf
-fisolate-erroneous-paths-dereference
-fira-remat
-foptimize-milling-calls
-fpartial-restrl
-fpartial-inlining
-fpeephole2
-freorder-blocks
-freorder-blocks-and-partition
-freorder-functions
-frerun-case-after-loop
-fschedule-interblock
-fsched-super
-fschedule-lsns
-fschedule-lsns2
-fstrict-aliasing
-fstrict-overflow
-ftree-built-in-call-dce
-ftree-switch-conversion
-ftree-tail-merge
-ftree-pre
-ftree-vrp
-fica-ra
-finline-functions
-funswitch-loops
-fpredictive-commoning
-fcase-after-reload
-ftree-loop-vectorize
-ftree-loop-distribute-patterns
-ftree-slp-vectorize
-fvect-cost-model
-ftree-partial-pce
-fipa-cp-clone

```

Parallelizing Compilers

Parallelization is just another optimization

- A compiler may automatically parallelize a sequential code
- Big quest during the 90s:
 - General parallelization is too hard for a compiler
 - Figuring out all data dependences is impossible at compile time
 - Relative success on more restrictive languages (CoArray Fortran, High Performance Fortran)
 - Vectorization of loops has become mainstream
- The programmer has to be involved in the parallelization of the code

How Do We Write Parallel Programs?

Parallel Programming Languages

During the 1990s

“C” in C	CUMMINS	Jazz BMI	P-RIO	Quake
ACE	DAPPLE	JAVA8	P4-Linda	Quick Threads
ACT++	Data Parallel C	Javaprocess	Pablo	Sage++
ADDAP	DC++	JDL	PADE	SAM
Ad	DC+++	Joyce	PADE	SCALOAL
Adsmith	IDL	Karma	Papers	Schedule
AFAPI	DRPC	Kleos	Para++	SeTIL
ALWAN	DSSM	KOAN/Fortran-S	Parc++	SDDA
AM	Distributed Smalltalk	LAM	Paradigm	SHMEM
AMG/C	DSM	Lam	Parasit	SIMPLEx
Ancoola	DOME	Linda	Parallelaci	Sisal
Appdall	DOSMOS	Linda	Parallel Haskell	SML
ARTS	DRL	LipSI	Parallel C++	SONIC
Athropacan-6b	DSM-Threads	Locut	Paro	Spago-C
Autonap	ECHO	Loud	ParLib++	SR
Autonap;	EIBAN	Mairie	ParLin	Streands
bdbthreads	Elisan	Massif	Parlog	Strand
Blaze	Emerald	Massifd	Parrot	SURE
BlockCணem	EPI	Meatat	Partrace	SWI
Blast	Metcalfe	Melange	ParTrig	SynergFascal
C*	Express	Milemp	pcC	Synergy
C**	Falcon	Millipede	pcL++	TCGMSG
C4	Filaments	Mirage	PCB	Telegraphone
CodeOS	FLASH	Model-2*	PCP	The FORTRAN
Cookware	FM	Modula-P	PCU	Threading++
CC++	Fork	MORIX	PRACE	TRAPPER
Charlotte	Fortran-M	MpC	PENNY	TreadMarks
Charm	FX	MPC++	PET	UC
Charm++	GAMMA	MPICH	PFTS	ufC++
Chi	GAMMA	Multipol	PH	UNITY
Cid	Gloria	Moxin	Phosphorus	V
CIS	GLU	Mnemo-Threads	POET	VIC*
CIS-Fortran	GUARD	NCI	POKE	Validifd V-NUS
Code	IMP	NetClassess++	POOL-T	VIPS
Concurrent ML	HORUS	Nexos	POOMA	Win32 threads
Converse	HPC	Nimrod	POSYBL	WinPar
COOL	HPC++	NOX	PRESTO	WWWWinds
CONVERGELATE	IMPACT	Objective Linda	Proteus	XENOOPS
CyberPar	IMPACT	Oasis	Protos	XPC
CPS	INET-Linda	Onaga	PSDM	Zornd
CRL	ISH	OOF90	PSI	ZPL
CSP	JADA	Orca	PVM	
Cthreads	JADE	P-e-e	QPC++	

Source: Patterns for Parallel Programming
(Tim Matson, et al, 2004)

- Based on **C/C++, Fortran, Python**

- A few dominant languages:

- Shared-memory: **OpenMP**
- Distributed-memory: **MPI**
- Accelerators: **OpenACC**

- Other programming models:

- Partitioned global address space:

UPC, OpenSHMEM

- Distributed arrays:

Charm++, GlobalArrays

- Task parallelism:

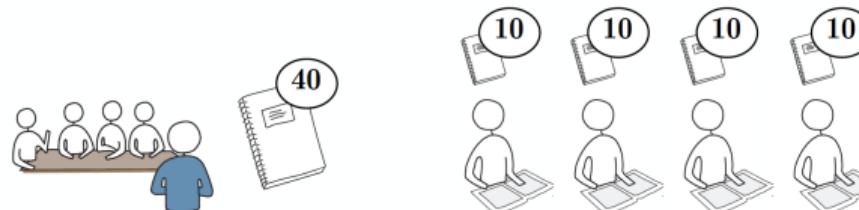
Uintah, Legion

- DARPA’s High Productivity Computer Systems:
Fortress, Chapel, X10

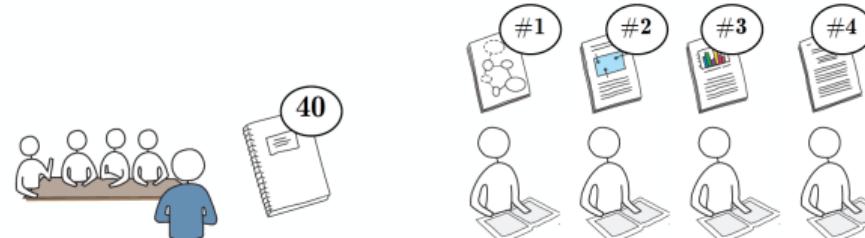
How Do We Write Parallel Programs?

Types of parallelism

- Data parallelism (SIMD): parallel execution of the same code stream on a supercomputer. Each node may process a different data portion.



- Task parallelism (MIMD): distribution of different code streams on a supercomputer. Each node may process a different data portion.



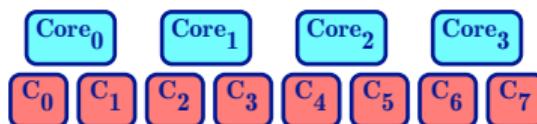
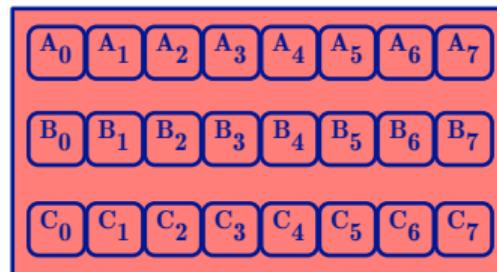
Parallel Programming

Considered harder than sequential programming

- Training in computer science is based on von Neumann architecture (sequential)
- Considerations in parallel programming:
 - Communication, sending data between processors
 - Synchronization, coordinating access to shared resources
 - Load balancing, distributing computation evenly across processor set
- Parallel machine:
 - Contains a set of p threads, numbered from 0 to $(p - 1)$
 - Each thread has a unique ID

Directive-based Parallel Programming

Add pragmas to hint compiler on parallel code



Sequential Code

```
1 for(i=0; i<N; i++)
2     C[i] = A[i] + B[i];
```

Definition (Pragmas)

Directive pragmas (from *pragmatics*) is a language construct to tell a compiler how it should process an input code. Pragmas are not part of the programming language.

Parallel Code

```
1 #pragma omp parallel
2 for(i=0; i<N; i++)
3     C[i] = A[i] + B[i];
```

Parallel Hello World

A primer on parallel execution

```
1 #include <omp.h>
2 int main(){
3     #pragma omp parallel
4     {
5         printf("Hello World\n");
6     }
7 }

1 #include <omp.h>
2 int main(){
3     int id, np;
4     #pragma omp parallel private(id, np)
5     {
6         np = omp_get_num_threads();
7         id = omp_get_thread_num();
8         printf("Hello from thread %d, out of %d threads\n", id, np);
9     }
10 }
```

Parallel Hello World

A primer on parallel execution

```
1 #include <omp.h>
2 int main(){
3     #pragma omp parallel
4     {
5         printf("Hello World\n");
6     }
7 }

1 #include <omp.h>
2 int main(){
3     int id, np;
4     #pragma omp parallel private(id, np)
5     {
6         np = omp_get_num_threads();
7         id = omp_get_thread_num();
8         printf("Hello from thread %d, out of %d threads\n", id, np);
9     }
10 }
```

Undo Parallelization

Finding a serial counterpart

- Directive-based parallel programming is based on the principle of *incremental parallelization*
- The *equivalent* serial program is always at hand:
 - To test for correctness
 - To compare performance
 - To get an more familiar serial algorithmic description

```
1 int main(){
2     #pragma omp parallel
3     BLOCK_A
4     ...
5     #pragma omp parallel
6     BLOCK_B
7     ...
8 }
```

```
1 int main(){
2     BLOCK_A
3     ...
4     BLOCK_B
5     ....
6 }
```

Exercises

- ① Download the git repository and inspect the openmp directory:

```
git clone https://github.com/CNCA-CeNAT/CRHPCS-2018.git
```

- ② Get into hello_id directory and get familiar with the execution mechanism:

- Type make to compile
- Type qsub hello_id.pbs to submit
- Open the output file to see the result

- ③ Get into map directory and implement a map program. Such program takes an array and applies a function foo to each member of the array. You must implement the foo function using any formula you want.

OpenMP

OpenMP

Open Multi-Processing

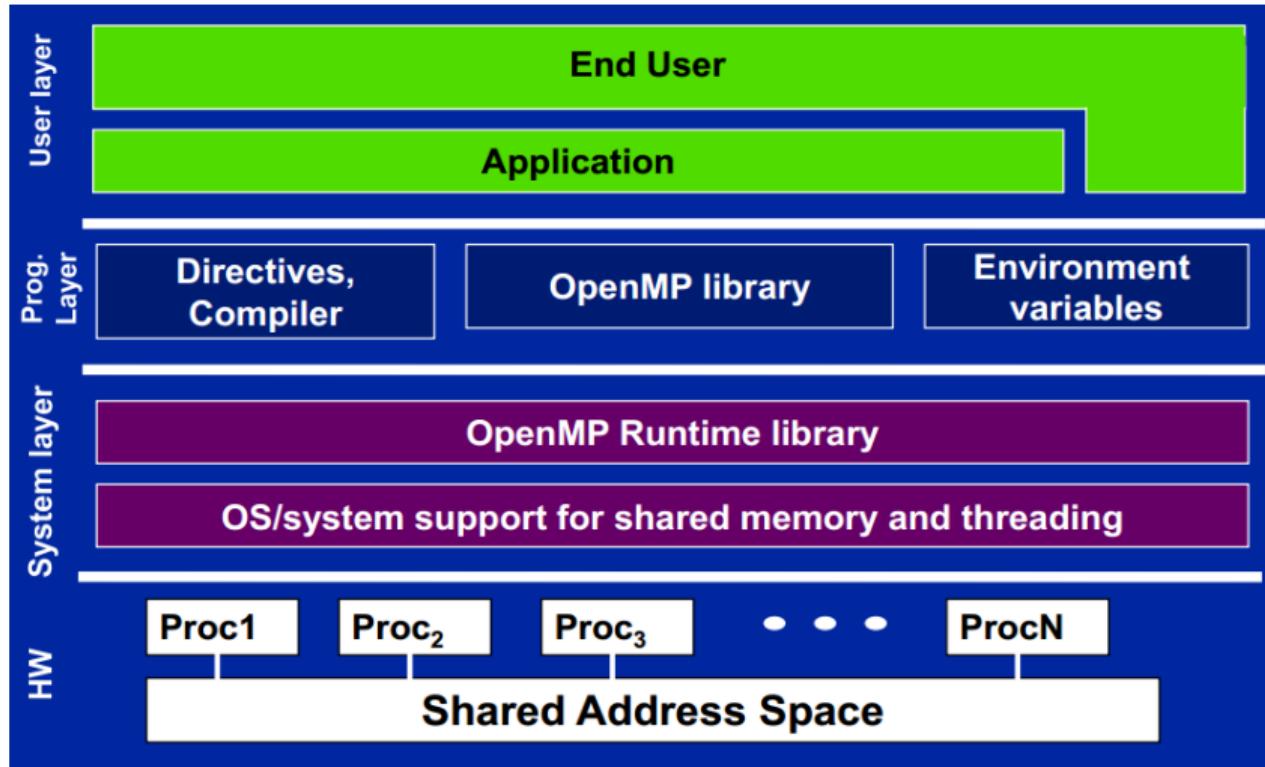
- An implementation of multithreading
- Standard API for shared-memory programming
- OpenMP Architecture Review Board (ARB) published OpenMP standard 1.0 in 1997
- Current release is OpenMP 4.5 in 2015
- Adopted by most important operating systems and supercomputer vendors
- Based on pragmas, programmer annotations for the compiler
- Embraces incremental parallelization philosophy
- Serial counterpart is always available

OpenMP Standard

Created by the community

- Defines an Application Program Interface (API) which provides a portable and scalable model for developers of shared memory parallel applications (instead of PThreads)
- A library with some simple functions
- A few program directives (pragmas = hints to the compiler)
- A few environment variables
- A compiler translates OpenMP functions and directives into PThread calls
- Program begins with a master thread
- Threads are forked when specified by directives

OpenMP Standard Solution Stack



High-level representation

Translated automatically into multithreaded code

OpenMP

```

1 #pragma omp parallel for
2 for(int i=0; i<n; i++){
3     A[i] = A[i] * A[i];
4 }
```

PThreads

```

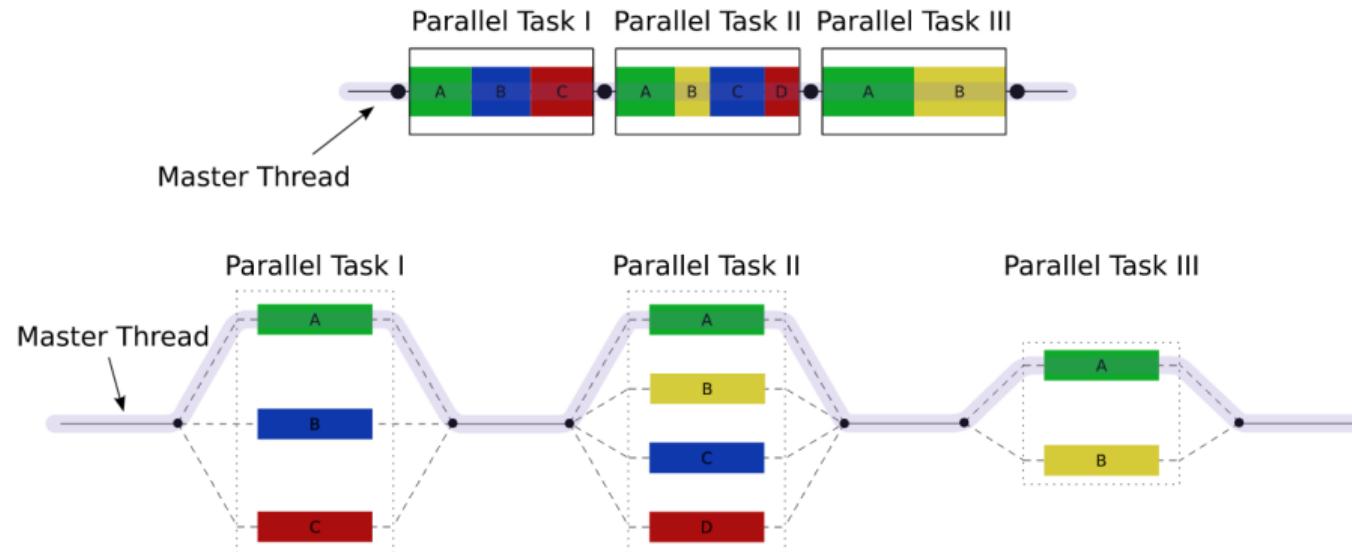
1 struct params {
2     int *A;
3     int tid;
4     int size;
5     int nthreads;
6 };
7 void *parallel_for(void *_p){
8     params *p          = (params*)_p;
9     int tid            = p->tid;
10    int chunk_size    = (p->size/p->nthreads);
11    int start          = tid * chunk_size;
12    int end            = start + chunk_size;
13    for(int i = start; i < end; i++)
14        A[i] = A[i] * A[i];
15    return 0;
16 }
```

```

17 pthread_t threads[nthreads];
18 params *thread_params = (params*) ? malloc(
19                                     nthreads * sizeof(params));
20 for(int i = 0; i < nthreads; i++){
21     thread_params[i].A      = A;
22     thread_params[i].tid    = i;
23     thread_params[i].size   = n;
24     thread_params[i].nthreads = nthreads;
25     pthread_create(&threads[i], NULL, ?
26                     parallel_for, (void*)
27                     &thread_params[i]);
28 }
29 for(int i = 0; i < nthreads; i++){
30     pthread_join(threads[i], NULL);
31 }
32 free(thread_params);
```

Fork-join Model

Parallelism shrinks and expands



Source: <http://www.wikipedia.org>

Parallel Regions

Defined with `#pragma omp parallel`

```
1 #pragma omp parallel
2 {
3     printf("Hello World\n");
4 }
```

The number of spawned threads can be specified through:

- Using the clause:

```
    num_threads(thread_count)
```

- An API function:

```
    void omp_set_num_threads(int number);
```

- An environment variable:

```
    export OMP_NUM_THREADS=32
```

```
1 #pragma omp parallel num_threads(32)
2 {
3     printf("Hello World\n");
4 }
```

```
1 omp_set_num_threads(32);
2 #pragma omp parallel
3 {
4     printf("Hello World\n");
5 }
```

Number of Threads

Define the level of concurrency

- The `num_threads(thread_count)` clause can be added to a parallel directive
- It allows the programmer to specify the number of threads that should execute the following block
- There may be system-defined limitations on the number of threads that a program can start
- The OpenMP standard does not guarantee that this will actually start `thread_count` threads
- Most current systems can start hundreds or even thousands of threads
- Unless we are trying to start a lot of threads, we will almost always get the desired number of threads

Embarrassingly Parallel

A problem that has a trivial parallel solution

- Simple distribution of work is enough to parallelize
- No significant communication complexity
- No significant synchronization complexity
- Should we feel embarrassed to come up with an easy parallel solution?
Of course not!
- Embarrassingly parallel → Delightfully parallel

Critical Sections and Barriers

Synchronizing threads

- A critical section is executed by one thread at a time; (name) is optional
- Critical sections with different names can be executed simultaneously

```
1 #pragma omp parallel
2 {
3     value = foo();
4     #pragma omp critical(name)
5     {
6         sum = sum + value;
7     }
8 }
```

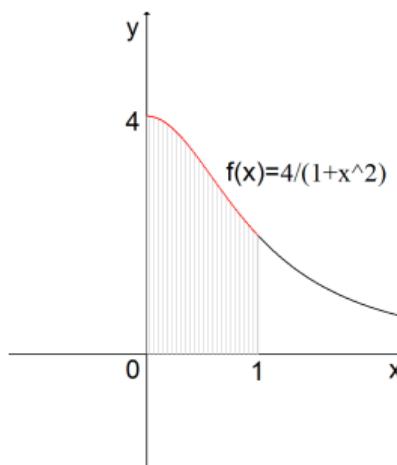
- All the threads in the active team must reach the barrier point before any can continue:

```
1 #pragma omp parallel
2 {
3     value = foo();
4     #pragma omp barrier
5     value = goo();
6 }
```

Exercise

Trapezoidal rule

- Approximating pi via numerical integration



- Divide interval up into subintervals.
- Assign subintervals to processes
- Each process calculates partial sum
- **Add all the partial sums together to get pi.**

Interacting with OpenMP Library

Useful functions and tips

- Compiler flag:

```
icc -fopenmp pgm.c
```

- Include file:

```
#include <omp.h>
```

- Environment variable for number of threads:

```
OMP_NUM_THREADS
```

- Useful functions:

- `omp_get_num_threads()`
returns the number of active threads
- `omp_get_thread_num()`
returns the thread ID
- `omp_set_num_threads(int t)`
sets number of active threads
- `omp_get_wtime()`
gets current wall clock time in seconds
- `omp_in_parallel()`
returns true if program is running in parallel; false otherwise

Data Scoping

Types of variables

- shared: accessible by all threads (default)
- private: one copy for each thread
- firstprivate: private but initialized to its value before entering the region
- lastprivate: private but on exit, value in master is last value in thread
- default: default type for all variables in the thread
- reduction: variables declared in enclosing context, are private in the parallel sections, but reduced upon exit using the specified operator (e.g.
+, *, -, &, |, ^, &&, ||)

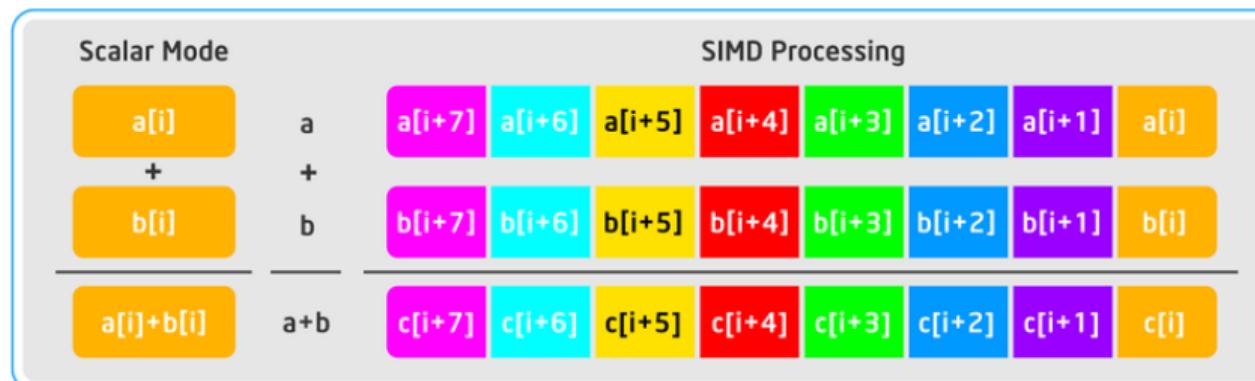
```
1 #pragma omp parallel private(id, np)
2 {
3     np = omp_get_num_threads();
4     id = omp_get_thread_num();
5     printf("Hello from thread %d, out of %d threads\n", id, np);
6 }
```

Vectorized Loops

Defined with `#pragma omp simd`

- Enforces vectorization of for loops
- The compiler may not automatically vectorize
- Several clauses and extensions

```
1 #pragma omp simd
2 for(i = 0; i < n; i++) {
3     A[i] = A[i] + K;
4 }
```



Parallel Loops

Defined with `#pragma omp for`

- Automatically distributes iterations in the for-loop to active threads

```
1 #pragma omp for
2 for(i=0; i<N; i++){
3     C[i] = A[i] + B[i];
4 }
```

- Usually used inside a parallel region

```
1 #pragma omp parallel
2 {
3     #pragma omp for
4     for(i=0; i<N; i++)
5         C[i] = A[i] + B[i];
6 }
```

- Pragmas `parallel` and `for` can be coalesced:

```
1 #pragma omp parallel for
2 for(i=0; i<N; i++)
3     C[i] = A[i] + B[i];
```

- Iteration index (variable `i` above) gets automatically privatized
- Reduction variables are useful often times

```
1 #pragma omp parallel for reduction
2 (+:sum)
3 for(int i=0; i<N; i++){
4     sum += vector[i];
5 }
```

Matrix Multiplication

Avoiding concurrency overhead

```
1 for (i=0; i<N; i++)
2     for (j=0; j<N; j++)
3         for (k=0; k<N; k++)
4             C[i][j] += A[i][k]*B[k][j];
```

Outer-loop parallelization

```
1 #pragma omp parallel for
2 for (i=0; i<N; i++)
3     for (j=0; j<N; j++)
4         for (k=0; k<N; k++)
5             C[i][j] += A[i][k]*B[k][j];
```

Inner-loop parallelization

```
1 for (i=0; i<N; i++)
2     #pragma omp parallel for
3     for (j=0; j<N; j++)
4         for (k=0; k<N; k++)
5             C[i][j] += A[i][k]*B[k][j];
```

Data Dependencies

Programmers must be careful to avoid race conditions

Fibonacci Sequence

$$f_n = f_{n-1} + f_{n-2}$$

Sequential Code

```
1 fibo[0] = fibo[1] = 1;  
2 for (i=2; i<10; i++)  
3     fibo[i] = fibo[i-1] + fibo[i-2];
```

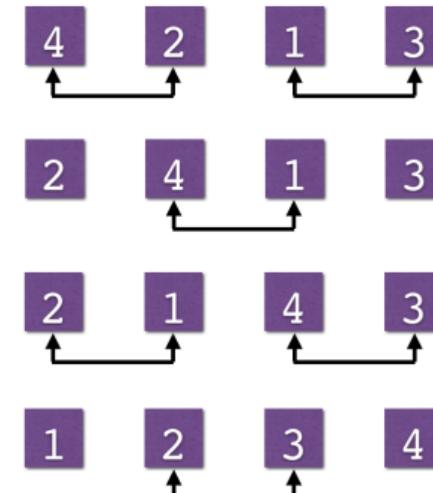
(Incorrect) Parallel Code

```
1 fibo[0] = fibo[1] = 1;  
2 omp_set_num_threads(2);  
3 #pragma omp parallel for  
4 for (i=2; i<10; i++)  
5     fibo[i] = fibo[i-1] + fibo[i-2];
```

Exercise

- ① Implement the odd-even parallel sort algorithm in the `odd_even_sort` directory. It is the parallel equivalent to *bubble sort*. When sorting n numbers with n processors, it alternates between $\frac{n}{2}$ odd and even phases:

- Each processor holds one number
- Odd phase: odd-numbered processors compare and swap their number with neighbor
- Even phase: even-numbered processors compare and swap their number with neighbor



Scheduling loops

Avoiding load imbalance

- Assignment of loops is system dependent
 - Normally block partitioning

```
1 sum = 0.0
2 for (i = 0; i <= n; i++)
3     sum += f(i)
```

- Good assignment of iterations to threads can have a very significant effect on performance.
- Assigning iterations to threads is called **scheduling**

Scheduling

Used to specify distribution of iterations to threads

- Syntax:

```
#pragma omp parallel for schedule(type,chunksize)
```

- Parameter type can be:

- *Static*: iterations are distributed among threads before the loop is executed
- *Dynamic or guided*: iterations are assigned to threads during the execution of the loop
- *Auto*: the compiler or runtime system determines the proper schedule
- *Runtime*: the environment variable OMP_SCHEDULE is used to determine schedule at runtime

- Parameter chunksize is a positive integer

Static Scheduling

Default distribution

- The iteration space is divided into chunks and distributed to threads in a round-robin fashion (default chunk size is 1)



- Example: 24 iterations with 4 threads

Chunksize	T1	T2	T3	T4
1	0,4,8,12,16,20	1,5,9,13,17,21	2,6,10,14,18,22	3,7,11,15,19,23
2	0,1,8,9,16,17	2,3,9,10,18,19	4,5,11,12,20,21	6,7,13,14,22,23
3	0,1,2,12,13,14	3,4,5,15,16,17	6,7,8,18,19,20	9,10,11,21,22,23
4	0,1,2,3,16,17,18,19	4,5,6,7,20,21,22,23	8,9,10,11	12,13,14,15

Dynamic Scheduling

Adapts to remaining work

- Iterations are divided into chunks and assigned to threads dynamically



- Each thread executes a chunk, and when done, requests another chunk from the runtime system
- Although each chunk contains the same number of iterations, chunks may have different execution times

Guided Scheduling

Big chunks first

- Similar to dynamic scheduling, except that chunk size decreases as the execution progresses



- At the dynamic scheduling decision, the chunk size is equal to $\frac{1}{p}$ of the remaining iterations, where p is the number of threads
- Can specify the smallest chunksize (except possibly the last)
- The default smallest chunksize is 1

Exercise

- 1 What would be the assignment of iterations to threads for a guided schedule in a parallel for loop with 1000 iterations and 2 threads?

The sequence in which threads 0 and 1 request more work is given by:

0,1,1,1,0,1,0,1,1,1,1.

Exercise

- ① What would be the assignment of iterations to threads for a guided schedule in a parallel for loop with 1000 iterations and 2 threads?

The sequence in which threads 0 and 1 request more work is given by:
0,1,1,1,0,1,0,1,1,1,1.

Chunk	Thread
[1-500]	0
[501-750]	1
[751-875]	1
[876-937]	1
[938-968]	0
[969-984]	1
[985-993]	0
[994-996]	1
[997-998]	1
999	1
1000	1

Sections

Pieces of codes that can run concurrently

- Define independent pieces of code
- May combine the parallel and sections pragmas (as with for pragma)
- If there are more threads than sections, then some are idle
- If there are fewer threads than sections, then some sections are serialized

```
1 #define N 1000
2 main(){
3     int i;
4     float a[N], b[N], c[N];
5     for (i=0; i < N; i++)
6         a[i] = b[i] = ... ;
7         #pragma omp parallel shared(a,b,c) private(i)
8         { ...
9             #pragma omp sections
10            {
11                #pragma omp section
12                {
13                    for (i=0; i < N/2; i++)
14                        c[i] = a[i] + b[i];
15                }
16                #pragma omp section
17                {
18                    for (i=N/2; i < N; i++)
19                        c[i] = a[i] + b[i];
20                }
21            } /* end of sections */
22            ...
23        } /* end of parallel section */
24    }
```

Single and Master

Define sections of code run by only one thread

- The single directive serializes a section of code within a parallel region
 - More convenient and efficient than terminating a parallel region and starting it later
 - Typically used to serialize a small section of the code that's not thread safe
 - Threads in the team that do not execute the single directive, wait at the end of the enclosed code block, unless a nowait clause is specified
- The master directive is the same as the single directive except that:
 - The serial code is executed by the master thread
 - The other threads skip the master section, but do not wait for the master thread to finish executing it

```
1 #pragma omp parallel num_threads(32)
2 {
3     printf("Hello World\n");
4     #pragma omp single
5     printf("Only one thread\n");
6     printf("Goodbye World\n");
7 }
```

Task Parallelism

Creating independent pieces of work

- Pragma task spawns a task (non-blocking)
- Pragma taskwait waits for all spawned tasks
- Pragmas for task parallelism must be used inside a parallel section
- **Taskloop construct:** loop iterations are partitioned into tasks and executed as tasks by the runtime system.

```
1 int fibo(int n){  
2     int fn1, fn2;  
3     if (n<2)  
4         return n;  
5     else {  
6         #pragma omp task  
7         fn1 = fibo(n-1);  
8         #pragma omp task  
9         fn2 = fibo(n-2);  
10        #pragma omp taskwait  
11        return fn1 + fn2;  
12    }  
13 }  
14 int main() {  
15     #pragma omp parallel shared(n)  
16     {  
17         #pragma omp single  
18         fibo(n);  
19     }  
20 }
```

Shared-memory Supercomputers

Powerful easy-to-program machines



- Cray XE1
- Multi streaming processors
- Up to 32 TB memory
- Up to 8192 processors



- SGI UV-2000
- Up to 64 TB memory
- Up to 256 processor sockets



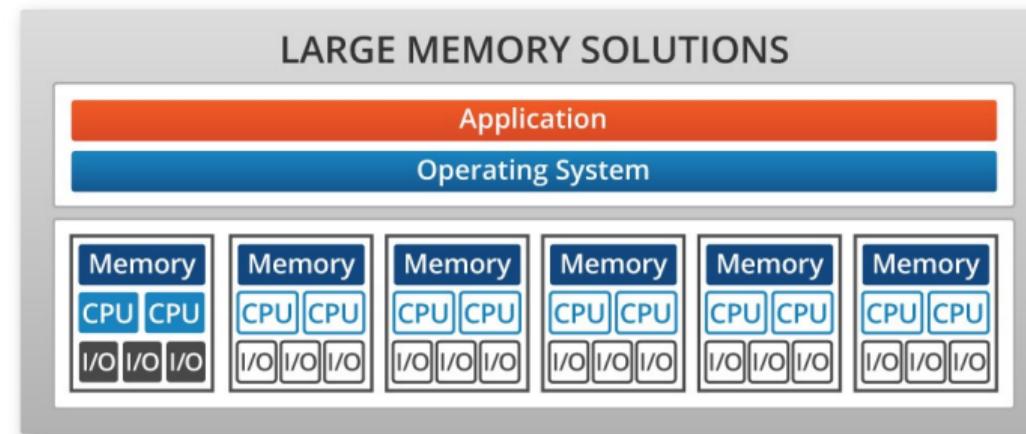
- Numaconnect
- Up to 256 TB memory
- Up to 4096 nodes

Virtual Shared Memory

Single address space via software



- Hypervisor or virtual machine monitor for HPC
- *Virtualization through aggregation*: combine multiple servers to create a symmetric multiprocessing system



Source: <http://www.hpcwire.com>

Using OpenMP

The Next Step

- The latest standard provides a way of controlling thread affinity (OpenMP threads placement)
 - Memory bandwidth
 - Latency
 - Cache utilization
 - First Touch Data Placement
- Heterogeneous Architectures
 - *Device Constructs*
 - GPUs, DSPs, FPGAs...
 - Threads cannot migrate across devices

Exercises

Exercise 1

Implement the parallel shear-sort algorithm.

Get into the `shear_sort` directory and read the description of the problem in file `shear_sort.pdf`.

Exercise 2

Implement the Strassen algorithm for parallel matrix multiplication.

Get into the `strassen` directory and read the description of the problem in file `strassen.pdf`.

Exercise 3

Think of any algorithm you may find relevant to your research interests and try to parallelize it using OpenMP pragmas.

Acknowledgements



Thank you!